

Project 3 - User-level Memory Management (100 points + Extra credit 10 points)

CS 416 - OS Design

Deadline: April 12th, 2021, 11:59pm

Please read the project description carefully.

Assume you are building a new startup for Cloud Infrastructure (Amazing Systems); a competitor of Amazon Cloud Services. As the CEO of your company, you decide to move page table and memory allocation to the software. In this project's part 1, you will build a user-level page table that can translate virtual addresses to physical addresses by building a multi-level page table. In part 2, you will also design and implement a translation lookaside buffer (TLB) cache to reduce translation cost. For evaluating your code, we will test your implementation across different page sizes. We will mainly focus on the correctness of your code, the page table and TLB state. For the extra-credit (see part 3), you will implement a 64-bit 4-level page table design. Note that this project would require understanding bit manipulation. We have added a helper code for you to learn and get used to bit manipulation (see 8. Sample Codes and Helpers).

Part 1. Implementation of Virtual Memory System (70 pts)

While you have used `malloc()` in the past, you might not have thought about how virtual pages are translated to physical pages and how they are managed. The goal of the project is to implement “`a_malloc()`” (Amazing malloc), which will return a virtual address that maps to a physical page. For simplicity, we will use a 32-bit address space that can support 4GB address space. **We will vary the physical memory size and the page size when testing your implementation.**

set_physical_mem(): This function is responsible for allocating memory buffer using `mmap` or `malloc` that creates an illusion of physical memory (Linux <http://man7.org/linux/man-pages/man2/mmap.2.htm>). Here, the physical memory refers to a large region of contiguous memory that is allocated using `mmap()` or `malloc()` and provides your page table and memory manager an illusion of physical memory. **Feel free to use malloc for allocating other data structures required for managing your virtual memory.**

translate(): This function takes input a page directory (address of the outer page table), a virtual address, and returns the corresponding physical address. You have to work with two-level page tables. For example, in a 4K page size configuration, each level uses 10-bits with 12-bits reserved for offset. For other page sizes X , reserve $\log_2(X)$ bits for offset and divide the remaining bits into two levels. Unequal division is possible in this case.

page_map(): This function walks the page directory to see if there is an existing mapping for a virtual address. If the virtual address is not present, then a new entry will be added. Note that you will be using this in `a_malloc()` to add page table entry.

a_malloc(): This function takes the number of bytes to allocate and returns a virtual address. Because you are using a 32-bit virtual address space, you are responsible for address management. To make things simple, assume for each allocation, you will allocate one or more pages (depending on the size of your allocation). For example, for `a_malloc(1byte)` and `a_malloc(1byte)`, you will allocate one page for each call. Similarly, for `a_malloc(4097 bytes)`, you will allocate 2 pages when the page size is 4KB. This simple allocation would cause fragmentation. One way to reduce *some* fragmentation can be found below:

Reducing Fragmentation (Optional-No Extra Credit): One way is to combine multiple small allocations to one page. For example, the first call of `a_malloc` returns an address 0x1000. When you call `a_malloc` again, you can return 0x1001 (if application asked for 1-byte) or any address within the page size boundary. The next call to `a_malloc` (if the size requested ends up not fitting within the first page) would return 0x2000 or higher. *Note: This is optional. Implement only if you have time. No extra points of this.*

Keeping Track: You will keep track of which physical pages are already allocated and which pages are free; use a virtual and physical page bitmap that represents a page (you **must** use a bitmap. Using any other structure will lose points, ie. using an array of `chars` with 'y' or 'n' or an array of ints using each index for a page). The `get_next_avail()` (see the code) function must return the next free available page. You must implement bitmaps efficiently (allocating only one bit per page) to avoid wasting memory (https://www.cprogramming.com/tutorial/bitwise_operators.html, also see the C Review [416 abridged] in resources on Sakai for some simple examples).

a_free(): This call takes a virtual address and the number of bytes (int), and releases (frees) pages starting from the page representing the virtual address. For example, for a 4K page size, `a_free(0x1000, 5000)` will free two pages starting at virtual addresses 0x1000. Also, please ensure `a_free()` isn't deallocating a page that hasn't been allocated yet! Note, `a_free` returns success only if all the pages are deallocated.

put_value(): This function takes a virtual address, a value pointer, and the size of the value pointer as an argument, and directly copies them to physical pages. Again, you have to check for the validity of the library's virtual address. Look into the code for implementation hints.

get_value(): This function also takes the same argument as `put_value()`, but reads the data from the virtual address to the value buffer. If you are implementing TLB, always check the presence of translation in TLB before proceeding forward.

mat_mult(): This function receives two matrices `mat1` and `mat2` as an argument with a size argument representing the number of rows and columns. After performing matrix multiplication, copy the result to answer array. Take a look at the test example. After reading the values from two matrices, you will perform multiplication and store them to an array. For indexing the matrices, use the following method:

```
A[i][j] = A[(i * size_of_rows * value_size) + (j * value_size)]
```

Important Notes: - You cannot change the function arguments/signature of the following: `a_malloc()`, `a_free()`, `put_value()`, `get_value()`, `mat_mult()`. - Your code must be thread-safe and your code will be tested with multi-threaded benchmarks.

We have included one sample test code (`benchmark/multi_test.c`) to check thread safety. This is just a sample, so please feel free to use, extend, increase thread count (and add to `Makefile`) the sample code for thread-safety checks.

Part 2. Implementation of a TLB (20 pts)

In this part, you will implement a direct-mapped TLB. Remember that a TLB caches virtual page number to physical address. This part cannot be completed unless Part 1 is correctly implemented.

Logic:

Initialize a direct-mapped TLB when initializing a page table. For any new page that gets allocated, no translation would exist in the TLB. So, after you add a new page table translation entry, also add a translation to the TLB by implementing `add_TLB()`.

Before performing a translation (in `translate()`), lookup the TLB to check if virtual to physical page translation exists. If the translation exists, you do not have to walk through the page table for performing translation (as done in Part 1). You must implement `check_TLB()` function to check the presence of a translation in the TLB.

If a translation does not exist in the TLB, check whether the page table has a translation (using your part 1). If a translation exists in the page table, then you could simply add a new virtual to physical page translation in the TLB using the function `add_TLB()`.

Number of entries in TLB: The number of entries in the TLB is defined in `my_vm.h` (`TLB_ENTRIES`). However, your code should work for any TLB entry count (modified using `TLB_ENTRIES` in `my_vm.h`).

TLB entry size: Remember, each entry in a TLB performs virtual to physical page translation. So, each TLB entry must be large enough to hold the virtual and physical page numbers.

TLB Eviction: As discussed in the class, the number of entries in a TLB are much lower than the number of page table entries. So clearly, we cannot cache all virtual to physical page translations in the TLB. Consequently, we must frequently evict some entries. A simple technique is to find the TLB index of a virtual page, and replace an older entry in the index with a new entry. The TLB eviction must be part of the `add_TLB()` function.

Expected Output: You must report the TLB miss rate by completing the `print_TLB_missrate()` function. See the class slides for the definition of TLB miss rate.

Important Note: Your code must be thread-safe and your code will be tested with multi-threaded benchmarks.

Part 3. Extra Credit: Implementing a 4-level page table (10 pts)

For the extra-credit, you are required to implement a 4-level page table for 64-bit addressing by extending the 2-level page table in Part 1 for 32-bit CPUs.

1. Before attempting to implement the extra-credit part, you must complete Part 1 and Part 2 of the project.
2. We cannot evaluate Part 3 unless Part 1 and 2 are implemented correctly.
3. Your 4-level design must be thread-safe.
4. Note that for the 64-bit design to work, you will have to compile the code without the `-m32` flag. Therefore, we suggest to create “`my_vm64.c`” and compile the code without the 32-bit flag.
5. Note that the TLB entries will also store 64-bit addresses.
6. For more questions, email us.

4. Suggested Steps

- Step 1. Design basic data structures for your memory management library.
- Step 2. Implement `set_physical_mem()`, `translate()` and `page_map()`. Make sure they work.
- Step 3. Implement `a_malloc()` and `a_free()`.
- Step 4. Test your code with matrix multiplication.
- Step 5. Implement a direct-mapped TLB if steps 1 to 4 work as expected.

5. Compiling and Benchmark Instructions

Please only use the given Makefiles for compiling. We have also provided a matrix multiplication benchmark to test the virtual memory functions. Before compiling the benchmark, you have to compile the project code first. Also, the benchmark does not display correct results until you implement your page table and memory management library. The benchmark provides hints for testing your code. Make sure your code is thread safe.

We will focus towards testing your implementation for correctness.

6. Report (10 pts)

Besides the VM library, you also need to write a report for your work. The report must include the following parts:

1. Detailed logic of how you implemented each virtual memory function.
2. Benchmark output for Part 1 and the observed TLB miss rate in Part 2.
3. Support for different page sizes (in multiples of 4K).
4. Possible issues in your code (if any).

5. If you implement the extra-credit part, description of the 4-level page table design and support for different page sizes.

Because we are using a 32-bit page table (for Part 1 and 2), the code compiles with -m32 flag. Not all iLab machines support -m32. Here's a list of them that you could use.

```
kill.cs.rutgers.edu
cp.cs.rutgers.edu
less.cs.rutgers.edu
ls.cs.rutgers.edu
```

7. Submission

Submit the following files to Sakai, as is (**Do not compress them**):

1. my_vm.h
2. my_vm.c
3. my_vm64.c (only if attempting extra-credit Part 3)
4. A report in .pdf format, completed, and with both partners names and netids on the report
5. Any other support source files and Makefiles you created or modified

Please Note: Your grade will be reduced if your submission does not follow the above instruction.

8. Sample Codes and Helpers.

This project requires understanding bit manipulation. In addition to the C tutorials we posted in the first week of the class, see a sample code (hw3-sample.c), which will help you understand bit manipulation.

In hw3-sample.c, you will see functions that show an example for the following operations. Remember, this is just a sample for you to understand the basics. You are responsible for understanding and trying out more complex bit manipulation (if required).

- 1) Extracting “N” top order bits (outer bits) from a 32-bit number
- 2) Extracting “M” middle-order bits from a 32-bit number
- 3) Setting “Kth bit” value in a character array
- 4) Reading “Kth bit” value in a character array

9. FAQs and Other Things To Note

1. Similar to project 2, we will update FAQs to a webpage, in addition to answering them in Piazza.

<https://www.cs.rutgers.edu/~sk2113/course/cs416-sp21/project3-faq.html>

2. **MAX_MEMSIZE vs. MEMSIZE.** Within the header file (*my_vm.h*), you will see two defined values: (1) MAX_MEMSIZE and (2) MEMSIZE. The difference between the two is that MAX_MEMSIZE is the size of the virtual address space you should support, while MEMSIZE is how much “physical memory” you should have. In this case MAX_MEMSIZE is defined as (4ULL * 1024 * 1024 * 1024) which is 2^{32} bytes or 4GB, the amount of virtual memory that a 32-bit system can have. On the other hand, MEMSIZE is defined as (1024 * 1024 * 1024) or 2^{30} bytes or 1GB, which is how much memory you should mmap() or malloc() to serve as your “physical memory”.
3. **Be mindful of values 2^{32} and up.** Notice that the definition of MAX_MEMSIZE is casted as an unsigned long long. This is because the library is compiled as a 32-bit program. With a 32-bit architecture, an int/unsigned int/long/unsigned long are all 4 bytes, meaning they can only represent the 0 to $2^{32} - 1$ different values. So when dealing with MAX_MEMSIZE or other values that equal or larger than 2^{32} , make sure to use unsigned long long to avoid any value truncation.

4. If you are using your personal computer for development and getting the following error, then refer to this link: <https://www.cyberciti.biz/faq/debian-ubuntu-linux-gnustubs-32-h-no-such-file-or-directory/gnu/stubs-32.h>: No such file or directory compilation terminated. make: *** ...